

Sparse Graph Label Randomization

October 19, 2023

1 Preliminaries

1.1 Bounded Functional Encryption

We will use the notation of static, bounded functional encryption as presented in [GGLW22].

Security

We will slightly weaken the security notion such that the adversary does not choose which circuits it can learn the functional secret key for. Indeed, this is a weaker notion of functional encryption which fixes the adversary's output circuit. We will assume that we get circuit C_1, \dots, C_d .

For completeness, we have the original security definition of [GGLW22] below:

$$\left\{ \begin{array}{l} (1^n, 1^q) \leftarrow \mathcal{A}^{(1)} \\ (\text{MPK}, \text{MSK}) \leftarrow \text{Setup}(1^n, 1^q) \\ m \leftarrow \mathcal{A}^{\text{KeyGen}(\text{MSK}, \cdot)}(\text{CT}) \\ \text{CT} \leftarrow \text{Enc}(\text{MPK}, m) \end{array} \right\}_{\lambda \in \mathbb{N}} \approx^c \left\{ \begin{array}{l} (1^n, 1^q) \leftarrow \mathcal{A}(1^\lambda) \\ (\text{MPK}, \text{st}_0) \leftarrow \text{Sim}_0(1^\lambda, 1^n, q) \\ m \leftarrow \mathcal{A}^{S_1(\text{st}_0)}(\text{MPK}) \\ (\text{CT}, \text{st}_2) \leftarrow \text{Sim}_2(\text{st}_1, \Pi^m) \end{array} \right\}_{\lambda \in \mathbb{N}}$$

whenever the following admissibility constraints and properties are satisfied:

- $\text{Sim}_1, \text{Sim}_3$ are stateful in that after each invocation, they updated their states st_1, st_3 respectively which is carried over to the next invocation.
- Π^m contains a list of functions f_i queried by \mathcal{A} in the pre-challenge phase along with their output on the challenge message m . That is, if f_i is the i -th function queried by \mathcal{A} to oracle Sim_1 and q_{pre} be the number of queries \mathcal{A} makes before outputting m , then $\Pi^m = ((f_1, f_1(m)), \dots, (f_{q_{\text{pre}}}, f_{q_{\text{pre}}}(m)))$.
- \mathcal{A} makes at most q queries combined to key generation oracle in both games.
- Sim_3 for each queried function f_i , in the post challenge phase, makes a single query to its message oracle U_m on the same f_i itself.

Our modified security definition is as follows:

$$\left\{ \begin{array}{l} (1^n, 1^q) \leftarrow \mathcal{A}^{(1)} \\ (\text{MPK}, \text{MSK}) \leftarrow \text{Setup}(1^n, 1^q) \\ m \leftarrow \mathcal{A}(\text{MPK}, \text{SK}_{C_1}, \dots, \text{SK}_{C_d}) \\ \text{CT} \leftarrow \text{Enc}(\text{MPK}, m) \end{array} \right\}_{\lambda \in \mathbb{N}} \approx^c \left\{ \begin{array}{l} (1^n, 1^q) \leftarrow \mathcal{A}(1^\lambda) \\ (\text{MPK}, \text{st}_0) \leftarrow \text{Sim}_0(1^\lambda, 1^n, q) \\ m \leftarrow \mathcal{A}^{S_1(\text{st}_0)}(\text{MPK}, C_1, \dots, C_d) \\ (\text{CT}, \text{st}_2) \leftarrow \text{Sim}_2(\text{st}_1, \Pi^m) \end{array} \right\}_{\lambda \in \mathbb{N}} \quad (1)$$

where the admissibility constraints remain the same.

1.2 Non-malleable Bounded FE

Here, we introduce the notion of non-malleable bounded functional encryption.

We define non-malleable security of bounded functional encryption in almost the exact notion of [Pas06] for public key encryption. First, let $NM(m_1, \dots, m_q, \mathcal{A})$ be a game as follows for $q = \text{poly}(\lambda)$:

1. $(\text{MPK}, \text{MSK}) \leftarrow \text{FE.Setup}(1^\lambda)$
2. $\text{CT}_1, \dots, \text{CT}_q \leftarrow \text{FE.Enc}(\text{MPK}, m_1), \dots, \text{FE.Enc}(\text{MPK}, m_q)$
3. $\text{CT}'_1, \dots, \text{CT}'_\ell \leftarrow \mathcal{A}(\text{MPK}, \text{CT}_1, \dots, \text{CT}_q, 1^{|m|})$
4. $m'_i \leftarrow \perp$ is $\text{CT}_i = \text{CT}'_j$ for any $i \in [q]$, $j \in [\ell]$ and $\text{FE.Dec}(\text{SK}_{\text{identity}}, c_i)$ otherwise.

Then, we say that a bounded functional encryption scheme is non-malleable if for all PPT \mathcal{A} and every PPT \mathcal{D} , there exists a negligible function negl such that for all $\{m\}_0, \{m\}_1 \in \{0, 1\}^{nq}$, we have

$$\left| \Pr[\mathcal{D}(NM(\{m\}_0, \mathcal{A})) = 1] - \Pr[\mathcal{D}(NM(\{m\}_1, \mathcal{A})) = 1] \right| \leq \text{negl}. \quad (2)$$

As outlined in [Pas06], we can equivalently define non-malleability in terms of a PPT recognizable relation R such that

$$\left| \Pr \left[NM(m_1, \dots, m_q, \mathcal{A}(z)) \in \bigcup_{m \in \{m\}} R(m) \right] - \Pr \left[c \leftarrow \text{Sim}_{NM}(1^n, z); m' = \text{FE.Dec}(\text{SK}_{\text{identity}}, c); m' \in \bigcup_{m \in \{m\}} R(m) \right] \right| \leq \text{negl}(\lambda). \quad (3)$$

Note that in the above definition, we do not give the adversary access to any SK_{C_i} . We simply require that the scheme is public key (many message) non-malleable.

2 Using Weak Extractible Obfuscation

2.1 Graph Randomized Traversal

Say that we have a sparse, potentially exponentially sized, graph $\mathcal{G} = (V, E)$ and $\forall v \in V, \deg(v) = d$. We also require that \mathcal{G} is equipped with a neighbor function, Γ , which can be computed in polynomial time. We define a randomized and keyed labelling function $\phi : \{0, 1\}^\lambda \times V \rightarrow \{0, 1\}^{\text{poly}(\lambda)}$ such that given, $\phi(K, v_0)$ for root v_0 , an adversary, \mathcal{A} , which does not know a path from v_0 to v ,

$$\Pr[\mathcal{A}(\mathcal{O}(C_\Gamma), v_0, v, \phi(K, v_0)) = \phi(K, v)] \leq \text{negl}(\lambda) \quad (4)$$

for function C_Γ where $C_\Gamma(\phi(K, u)) = \phi(K, \Gamma(u)_1), \dots, \phi(K, \Gamma(u)_d)$ if $\Gamma(u) \neq \emptyset$ and otherwise $\Gamma(u)$ returns a \perp string; and, \mathcal{O} represents an indistinguishable obfuscator.

2.2 Instantiation

We define

$$\phi(K, v) = F(K, v).$$

For shorthand, we will write σ_v to connote an attempted “signature” of v where a correct signature is $F(K, v)$.

We can now define C_Γ :

Algorithm 1 The circuit for the neighbor function, C_Γ .

```

1: function  $C_\Gamma(f(\sigma_v), v)$ 
2:   if  $f(\sigma_v) \neq f(F(K, v))$  then
3:     return  $\perp$ 
4:   if  $\Gamma(v) = \emptyset$  then
5:     return  $\perp$ 
6:    $u_1, \dots, u_d = \Gamma(v)$ 
7:   return  $f(F(K, u_1)), f(F(K, u_2)), \dots, f(F(K, u_d))$ 

```

We are going to show that [eq. \(4\)](#) holds by first showing that the non-existence of an extractor to find a path from v_0 to v implies that \mathcal{A} necessarily does not know $\phi(K, c)$ for a $c \in C_V \subset V$ where the vertices in C_V border a graph cut which separates v_0 and v . Then, we inductively build up a series of games to show that \mathcal{A} cannot learn *any* $\phi(K, v)$ for $v \in V_1$ where V_1 are the vertices on the right-hand side of the cut.

Lemma 2.1. *Assuming that there is no extractor E such that $\Pr[E(\Gamma, v_0, v) = P] \geq \frac{1}{p(\lambda)}$ where $P \in \mathcal{P}$, then for any PPT \mathcal{A} , there exists some graph cut $C_E \subset E$ which separates v_0 and v and a set C_V such that*

$$\Pr[\mathcal{A}(\mathcal{O}(C_\Gamma), v_0, v, \phi(K, v_0)) \in \phi(K, C_V)] \leq \text{negl}(\lambda). \quad (5)$$

We define $C_V \subset V$ to be

$$\{u \mid (w, u) \in C_E \text{ and } u \text{ on the side of } v\} \cup \{v \mid (w, u) \in C_E \text{ and } u \text{ on the side of } v\}.$$

In words, C_V are the vertices just adjacent to the cut and on the same side as v .

Proof. We will show that if \mathcal{A} can break [eq. \(5\)](#), then we can construct an extractor, E , which finds a path from v_0 to v with non-negligible probability.

Assume that for every possible cut, \mathcal{A} is able to produce a single label in this cut for a vertex w . Then, we note that there must be at least 1 path from v_0 to w and v as otherwise, w would not be in the cut. Moreover, we note that \mathcal{A} must be able to produce a label for all vertices on at least one path from v_0 to w as otherwise, we can change the cut to include the edges between where \mathcal{A} is able to produce a label and not able to produce a label. Using the same argument, we can show that \mathcal{A} must be able to produce all labels on a path from w to v .

Note that \mathcal{A} is not given the specific cut C_E but rather C_E is chosen based off of the adversary. So, we can build an extractor to do the following:

1. Create an iO obfuscated circuit with a random key, K' , for C_Γ and create circuit $\mathcal{O}(C_\Gamma)$ as well as $\phi(K', v_0)$
2. Run $\mathcal{A}(\mathcal{O}(C_\Gamma), v_0, v, \phi(K', v_0))$ to get all labels $\phi(K', v_0), \dots, \phi(K', v)$ for some path from v_0 to v .
3. Recreate the path from v_0 to v via checking which vertex matches to adjacent labels in the path: I.e. starting with $\ell = 0$, we can learn the $\ell + 1$ vertex via finding $j \in [d]$ such that $C_\Gamma(\phi(K', v_\ell), v_\ell)_j \in \{\phi(K', v_0), \dots, \phi(K', v)\}$ and then setting $v_{\ell+1} = \Gamma(v_\ell)_j$.

□

Algorithm 2 Circuit for the neighbor function, $C_\Gamma^{w^*, \Gamma(w^*)_1, \dots, \Gamma(w^*)_d}$ with punctured PRF key $K(\{w^*\})$ and constant $z^*, z_1^*, z_2^*, \dots, z_d^*$

```

1: function  $C_\Gamma(f(\sigma_v), v)$ 
2:   if  $v \neq w$  and  $f(\sigma_v) \neq f(F(K, v))$  then
3:     return  $\perp$ 
4:   if  $v = w$  and  $f(\sigma_v) \neq z^*$  then
5:     return  $\perp$ 
6:   if  $\Gamma(v) = \emptyset$  then
7:     return  $\perp$ 
8:   if  $v = w$  then
9:     return  $z_1^*, z_2^*, \dots, z_d^*$ 
10:   $u_1, \dots, u_d = \Gamma(v)$ 
11:  return  $f(F(K, u_1)), f(F(K, u_2)), \dots, f(F(K, u_d))$ 

```

We are going to use a series of inductively built indistinguishable hybrids along with [algorithm 2](#) to show that [eq. \(4\)](#) holds.

- **Hyb₀**: In the first hybrid, the following game is played

1. $K \xleftarrow{\$} \{0, 1\}^{\lambda'}$ and $\phi(K, v_0) = (F(K, v_0), v)$
2. The challenger generates $\mathcal{O}(C_\Gamma)$ and gives the program to \mathcal{A}
3. The challenger chooses a v and gives the adversary v in plaintext.
4. \mathcal{A} outputs guess g and wins if $g = \phi(K, v)$

- **Hyb₁**: Let \mathcal{P} be the set of all paths from v_0 to v . For each path $P \in \mathcal{P}$ where P is an ordered list of connected vertices, we have that the adversary does not know some part of P . We can note that this implies that \mathcal{A} does not know $\phi(K, p)$ for all $p \in P$ as then \mathcal{A} can recover P . Let u_P be the first vertex in P such that \mathcal{A} does not know a path from v_0 to u_P . Define $\text{Suff}'(P)$ to be the path in P from this u_P to v . Then, necessarily, \mathcal{A} does not know $\phi(K, w_P)$ for at least one $w_P \in \text{Suff}'(P)$ as then \mathcal{A} would know a path from v_0 to v . Now, let $\text{Suff}(P)$ be the path which starts at w_P , ends at v .

We now inductively build up a series of hybrids to show that a hybrid distribution which shows $\phi(K, s)$ for $s \in \text{Suff}(P)$ indistinguishable from random. We perform the following procedure for each $P \in \mathcal{P}$. So, for $P \in \mathcal{P}$,

1. For the base case, let $U = \text{Suff}(P)_1$ where $\text{Suff}(P)_1$ is the first vertex in P such that \mathcal{A} does not know $\phi(K, p)$ for $p \in P$.
 - (a) Set $w^* = p$. Then, replace C_Γ with $C_\Gamma^{w^*, \Gamma(w^*)_1, \dots, \Gamma(w^*)_d}$ as defined in [algorithm 2](#). Fix the constant $z^* = f(F(K, p))$ and $z_1^* = f(F(K, \Gamma(w^*)_1)), \dots, z_d^* = f(F(K, \Gamma(w^*)_d))$.
 - (b) Set $z^* = f(t), z_1^* = f(t_1), \dots, z_d^* = f(t_d)$ where t, t_1, \dots, t_d are chosen at random
2. For the ℓ -th inductive step where $1 \leq \ell < |\text{Suff}(P)|$, we are going to assume that we are given a hybrid such that $w^* = \text{Suff}(P)_\ell$ and $z^* = f(t), z_1^* = f(t_1), \dots, z_d^* = f(t_d)$ for random t, \dots, t_d in [algorithm 2](#). Now, we change the hybrid in a similar manner as in the base case:
 - (a) Set $w^* = \text{Suff}(P)_{\ell+1}$. Then, replace C_Γ with $C_\Gamma^{w^*, \Gamma(w^*)_1, \dots, \Gamma(w^*)_d}$ as defined in [algorithm 2](#). Fix the constant $z^* = f(F(K, p))$ and $z_1^* = f(F(K, \Gamma(w^*)_1)), \dots, z_d^* = f(F(K, \Gamma(w^*)_d))$.
 - (b) Set $z^* = f(t), z_1^* = f(t_1), \dots, z_d^* = f(t_d)$ where t, t_1, \dots, t_d are chosen at random to puncture on $\text{Suff}(P)_{\ell+1}$ where we update z^*, \dots, z_d^* with new randomness.

Finally, we can note that if $\text{Hyb}_0 \stackrel{c}{\approx} \text{Hyb}_1$,

$$\Pr[\mathcal{A}(C_\Gamma, v_0, v, \phi(K, v_0)) = \phi(K, v)] \stackrel{c}{\approx} \Pr[\mathcal{A}(C_\Gamma^*, v_0, v, \phi(K, v_0)) = \phi(K, v)]$$

where z^* in C_Γ^* is the image on a OWF of a randomly chosen point. Thus, if \mathcal{A} can produce $\phi(K, v) = (\sigma_v, v)$, then \mathcal{A} can find a preimage for z^* under f and thus break the security of a one way function.

Now, we show that $\text{Hyb}_0 \stackrel{c}{\approx} \text{Hyb}_1$.

Lemma 2.2. *The game in $\text{Hyb}_1(1a)$ is indistinguishable from Hyb_0 .*

Proof. As the functionality of C_Γ in Hyb_0 equals that of $\text{Hyb}_1(1a)$, we have indistinguishable simply from the definition of indistinguishable obfuscation. \square

Lemma 2.3. *The game in $\text{Hyb}_1(1b)$ is indistinguishable from $\text{Hyb}_1(1a)$.*

Proof. Here we argue that if the game in $\text{Hyb}_1(1b)$ is distinguishable from $\text{Hyb}_1(1a)$, then we can construct an adversary, \mathcal{B} , which can break the security of the PRF at the punctured point. \square

Lemma 2.4. *The game in $\text{Hyb}_1(2a)$ is indistinguishable from Hyb_0 and, by the inductive hypothesis, all previous hybrids.*

Proof. Again, we have that the circuit for C_Γ is the same in Hyb_0 and $\text{Hyb}_1(2a)$. Thus, by the definition of indistinguishable obfuscation, these games are indistinguishable. \square

Lemma 2.5. *The game in $\text{Hyb}_1(2b)$ is indistinguishable from $\text{Hyb}_1(2a)$ and, by the inductive hypothesis, all previous hybrids.*

Proof. TODO: PRF security + extractor part \square

October 19, 2023

Abstract

References

- [GGLW22] Rachit Garg, Rishab Goyal, George Lu, and Brent Waters. Dynamic collusion bounded functional encryption from identity-based encryption. In *Annual International Conference on the Theory and Applications of Cryptographic Techniques*, pages 736–763. Springer, 2022. [1.1](#), [1.1](#)
- [Pas06] Rafael Pass. Lecture 16: Non-malleability and public key encryption, October 2006. [1.2](#), [1.2](#)